

Rain Fall disrict wise dataset

Problem statement: Analysing the monthly rainfall data and comparing with Annual data.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df=pd.read_csv(r"C:\Users\yoshitha lakshmi\OneDrive\Desktop\python\Rainfall district wise.csv")
        2 df
```

Out[2]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.9	354.8	326.0	315.2	250.9	2805.2	165.2	540.0
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.1	455.6	301.2	275.8	128.3	3015.7	69.7	480.0
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.9	454.8	276.1	198.6	100.0	2913.3	48.6	400.0
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.8	313.6	167.1	34.1	29.8	3043.8	123.0	840.0
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.2	568.0	206.9	29.5	31.7	4034.7	112.8	640.0
...
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527.3	308.4	343.2	172.9	48.1	3302.5	35.5	420.0
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636.3	263.1	234.9	84.6	18.4	3621.6	3.3	270.0
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352.7	266.2	359.4	213.5	51.3	2958.4	65.0	550.0
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592.9	230.7	213.1	93.6	25.8	3253.1	13.1	270.0
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217.5	163.1	157.1	117.7	58.8	1600.0	35.5	230.0

641 rows × 19 columns



In [3]:

1 df.head()

Out[3]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.9	354.8	326.0	315.2	250.9	2805.2	165.2	540.7	1207.2
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.1	455.6	301.2	275.8	128.3	3015.7	69.7	483.5	1757.2
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.9	454.8	276.1	198.6	100.0	2913.3	48.6	405.6	1884.4
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.8	313.6	167.1	34.1	29.8	3043.8	123.0	841.3	1848.5
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.2	568.0	206.9	29.5	31.7	4034.7	112.8	645.4	3008.4

In [4]:

1 df.tail()

Out[4]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527.3	308.4	343.2	172.9	48.1	3302.5	35.5	426.6
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636.3	263.1	234.9	84.6	18.4	3621.6	3.3	272.9
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352.7	266.2	359.4	213.5	51.3	2958.4	65.0	553.5
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592.9	230.7	213.1	93.6	25.8	3253.1	13.1	275.4
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217.5	163.1	157.1	117.7	58.8	1600.0	35.5	232.4


In [5]: 1 df.shape

Out[5]: (641, 19)

In [6]: 1 df.describe()

Out[6]:

	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV
count	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000
mean	18.355070	20.984399	30.034789	45.543214	81.535101	196.007332	326.033697	291.152262	194.609048	90.446334	34.117473
std	21.082806	27.729596	45.451082	71.556279	111.960390	196.556284	221.364643	152.647325	99.830540	74.990685	59.371274
min	0.000000	0.000000	0.000000	0.000000	0.900000	3.800000	11.600000	14.100000	8.600000	3.100000	1.200000
25%	6.900000	7.000000	7.000000	5.000000	12.100000	68.800000	206.400000	194.600000	128.800000	34.300000	6.600000
50%	13.300000	12.300000	12.700000	15.100000	33.900000	131.900000	293.700000	284.800000	181.300000	62.600000	12.900000
75%	19.200000	24.100000	33.200000	48.300000	91.900000	226.600000	374.800000	358.100000	234.100000	130.200000	32.300000
max	144.500000	229.600000	367.900000	554.400000	733.700000	1476.200000	1820.900000	1522.100000	826.300000	517.700000	475.100000



In [7]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 641 entries, 0 to 640
Data columns (total 19 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   STATE_UT_NAME         641 non-null   object 
 1   DISTRICT              641 non-null   object 
 2   JAN                   641 non-null   float64
 3   FEB                   641 non-null   float64
 4   MAR                   641 non-null   float64
 5   APR                   641 non-null   float64
 6   MAY                   641 non-null   float64
 7   JUN                   641 non-null   float64
 8   JUL                   641 non-null   float64
 9   AUG                   641 non-null   float64
10  SEP                   641 non-null   float64
11  OCT                   641 non-null   float64
12  NOV                   641 non-null   float64
13  DEC                   641 non-null   float64
14  ANNUAL                641 non-null   float64
15  Jan-Feb               641 non-null   float64
16  Mar-May               641 non-null   float64
17  Jun-Sep               641 non-null   float64
18  Oct-Dec               641 non-null   float64
dtypes: float64(17), object(2)
memory usage: 95.3+ KB
```

In [15]:


```
1 features=df.columns[2:13]
```

In [16]:

```
1 target=df.columns[14]
```

```
In [17]: FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC'], y_vars=['ANNUAL'], height=7, aspect=0.7, kind='reg')
```

```
Out[17]: <seaborn.axisgrid.PairGrid at 0x1f097f97f40>
```



```
In [18]: 1 x=np.array(df[features])
          2 y=np.array(df[target])
```

```
In [19]: 1 from sklearn.model_selection import train_test_split
          2 from sklearn.linear_model import LinearRegression
          3 x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.25)
          4 lm=LinearRegression()
          5 lm.fit(x_train,y_train)
          6 print(lm.score(x_train,y_train))
```

```
0.9998981063634227
```

```
In [20]: 1 coeff_df=pd.DataFrame(lm.coef_)
          2 coeff_df
```

Out[20]:

	0
0	1.547951
1	0.594910
2	1.233032
3	1.011594
4	0.962484
5	1.005697
6	1.001794
7	0.978258
8	1.065904
9	0.795722
10	1.644514

```
In [36]: 1 predictions=lm.predict(x_test)
```

```
In [25]: 1 from sklearn import metrics
          2 print('MAE:',metrics.mean_absolute_error(y_test,predictions))
          3 print('MSE:',metrics.mean_squared_error(y_test,predictions))
          4 print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```


MAE: 5.915817136311949
MSE: 75.99669382142075
RMSE: 8.717608262672782

Exploratory Data Analysis

```
In [27]: 1 raaindf=df[['JAN','FEB','MAR','APR','MAY','JUN','JUL','AUG','SEP','OCT','NOV','DEC','ANNUAL']]
```


```
In [28]: 1 sns.heatmap(raaindf.corr())
```

Out[28]: <Axes: >





```
In [39]: 1 predictions=lm.predict(x_test)
          2 plt.scatter(y_test,predictions)
```

```
Out[39]: <matplotlib.collections.PathCollection at 0x1f0b601a620>
```




```
In [43]: 1 sns.displot((y_test-predictions),bins=50);
```



```
In [44]: 1 sns.displot(df[ 'ANNUAL ' ])
```

```
Out[44]: <seaborn.axisgrid.FacetGrid at 0x1f0b5da4c10>
```



Ridge Regression Model


```
In [46]: 1 from sklearn.linear_model import Ridge,Lasso,RidgeCV
2 ridgeReg = Ridge(alpha=10)
3 ridgeReg.fit(x_train,y_train)
4 train_score_ridge = ridgeReg.score(x_train,y_train)
5 test_score_ridge = ridgeReg.score(x_test,y_test)
6
7 print('\nRidge Model\n')
8 print('Train score for ridge model is {}'.format(train_score_ridge))
9 print('Test score for ridge model is {}'.format(test_score_ridge))
```

Ridge Model

Train score for ridge model is 0.9998981062781498

Test score for ridge model is 0.9998869925352645

```
In [52]: 1 plt.figure(figsize = (10,4))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Ridge;$\alpha=10$')
3 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```




```
In [56]: 1 lassoReg = Lasso(alpha=10)
          2 lassoReg.fit(x_train,y_train)
          3 train_score_lasso = lassoReg.score(x_train,y_train)
          4 test_score_lasso = lassoReg.score(x_test,y_test)
          5
          6 print('\nRidge Model\n')
          7 print('Train score for lasso model is {}'.format(train_score_lasso))
          8 print('Test score for lasso model is{}'.format(test_score_lasso))
```


Ridge Model

Train score for lasso model is 0.9998976095395863

Test score for lasso model is0.999882011015853


```
In [57]: 1 plt.figure(figsize = (10,10))
2 plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Lasso;$\lambda$')
3 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regres')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```







In [58]:

```
1 # Comparison between Ridge,Lasso and RidgeCV
2 plt.figure(figsize=(10,10))
3
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Ridge;$\
5 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='d',markersize=6,color='red',label=r'Lasso;$\
6 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regres
7 plt.xticks(rotation=90)
8 plt.title('comparison between ridge,lasso and linear')
9 plt.legend()
10 plt.show()
```



```
In [62]: 1 # Linear CV model using Ridge
2 from sklearn.linear_model import RidgeCV
3 ridge_CV=RidgeCV(alphas=[0.0001,0.01,0.001,0.1]).fit(x_train,y_train)
4 print("The train score for ridge model is {}".format(ridge_CV.score(x_train,y_train)))
5 print("The test score for ridge model is {}".format(ridge_CV.score(x_test,y_test)))
```

The train score for ridge model is 0.9998981036461831

The test score for ridge model is 0.9998869496121741

```
In [63]: 1 # Linear CV model using Lasso
2 from sklearn.linear_model import LassoCV
3 lasso_CV=LassoCV(alphas=[0.0001,0.01,0.001,0.1]).fit(x_train,y_train)
4 print("The train score for lasso model is {}".format(lasso_CV.score(x_train,y_train)))
5 print("The test score for lasso model is {}".format(lasso_CV.score(x_test,y_test)))
```

The train score for lasso model is 0.9998980995382691

The test score for lasso model is 0.9998872120915944

```
In [60]: 1 from sklearn.linear_model import ElasticNet
2 regr = ElasticNet()
3 regr.fit(x,y)
4 print(regr.coef_)
5 print(regr.intercept_)
6 y_pred_elastic = regr.predict(x_train)
7 mean_squared_error = np.mean((y_pred_elastic-y_train)**2)
8 print('Mean squared error on test set',mean_squared_error)
```

[1.59402227 0.77496303 1.07666284 1.00645481 0.98815802 1.00582305
0.99931781 0.97704429 1.06450887 0.81537098 1.62088766]
-0.9346307788046033
Mean squared error on test set 88.18210015402676

```
In [65]: 1 regr.score(x_train,y_train)
```

Out[65]: 0.999888082161808

Conclusion

The Accuracy for LinearRegression model is 0.99 i.e, the data is fitted well where the accuracy for Ridge and Lasso and Elastic Net is also same.

Rainfall 1901-2015 dataset

Problem statement: Analysed the annual rainfall data for different years.

```
In [66]: 1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

In [67]:

```
1 df=pd.read_csv(r"C:\Users\yoshitha lakshmi\OneDrive\Desktop\python\Rainfall 1901-2015.csv")
2 df
```

Out[67]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	C
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3	98
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.9	71
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.0	69
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.6	57
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9	63
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	1013.0	31
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	1119.5	16
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	1057.0	17
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	958.5	29
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	860.9	55

4116 rows × 19 columns



In [68]:

1 df.head()

Out[68]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3	980.3
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.9	716.7
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.0	690.6
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.6	571.0
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9	630.8

In [69]:

1 df.tail()

Out[69]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	1013.0	316.6
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	1119.5	167.1
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	1057.0	177.6
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	958.5	290.5
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	860.9	555.4



In [70]:

```
1 df.describe()
```

Out[70]:

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	
count	4116.000000	4112.000000	4113.000000	4110.000000	4112.000000	4113.000000	4111.000000	4109.000000	4112.000000	4110.000000	4109.
mean	1958.218659	18.957320	21.805325	27.359197	43.127432	85.745417	230.234444	347.214334	290.263497	197.361922	95.
std	33.140898	33.585371	35.909488	46.959424	67.831168	123.234904	234.710758	269.539667	188.770477	135.408345	99.
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000	0.000000	0.000000	0.100000	0.
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.350000	175.600000	155.975000	100.525000	14.
50%	1958.000000	6.000000	6.700000	7.800000	15.700000	36.600000	138.700000	284.800000	259.400000	173.900000	65.
75%	1987.000000	22.200000	26.800000	31.300000	49.950000	97.200000	305.150000	418.400000	377.800000	265.800000	148.
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000	2362.800000	1664.600000	1222.000000	948.

In [71]:

```
1 df.shape
```

Out[71]: (4116, 19)

In [72]:

```
1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   SUBDIVISION     4116 non-null   object  
 1   YEAR            4116 non-null   int64   
 2   JAN             4112 non-null   float64  
 3   FEB             4113 non-null   float64  
 4   MAR             4110 non-null   float64  
 5   APR             4112 non-null   float64  
 6   MAY             4113 non-null   float64  
 7   JUN             4111 non-null   float64  
 8   JUL             4109 non-null   float64  
 9   AUG             4112 non-null   float64  
10  SEP             4110 non-null   float64  
11  OCT             4109 non-null   float64  
12  NOV             4105 non-null   float64  
13  DEC             4106 non-null   float64  
14  ANNUAL          4090 non-null   float64  
15  Jan-Feb         4110 non-null   float64  
16  Mar-May         4107 non-null   float64  
17  Jun-Sep         4106 non-null   float64  
18  Oct-Dec         4103 non-null   float64  
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [93]:

```
1 df.fillna(method='ffill', inplace=True)
```

In [94]:

```
1 features=df.columns[1:13]
```


```
In [95]: 1 target=df.columns[14]
```

```
In [96]: 1 df.columns
```

```
Out[96]: Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',  
              'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May',  
              'Jun-Sep', 'Oct-Dec'],  
             dtype='object')
```

```
In [97]: 1 sns.pairplot(df,x_vars=['YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',  
2          'AUG', 'SEP', 'OCT', 'NOV', 'DEC'],y_vars=['ANNUAL'],height=7,aspect=0.5,kind='reg')
```

```
Out[97]: <seaborn.axisgrid.PairGrid at 0x1f0be4e8940>
```



```
In [98]: 1 x=np.array(df[features])  
2 y=np.array(df[target])
```

```
In [99]: 1 from sklearn.model_selection import train_test_split  
2 from sklearn.linear_model import LinearRegression  
3 x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.25)  
4 lm=LinearRegression()  
5 lm.fit(x_train,y_train)  
6 print(lm.score(x_train,y_train))
```

```
0.9916135364757804
```

```
In [102]: 1 coeff_df=pd.DataFrame(lm.coef_)
          2 coeff_df
```

Out[102]:

	0
0	-0.036257
1	1.094004
2	0.792353
3	1.285093
4	1.018243
5	0.956323
6	0.995912
7	1.017815
8	0.917773
9	1.033492
10	1.061315
11	1.282726

```
In [103]: 1 print(lm.intercept_)
```

82.37961923825105

```
In [104]: 1 predictions=lm.predict(x_test)
```

In [105]:

```
1 from sklearn import metrics
2 print('MAE:', metrics.mean_absolute_error(y_test, predictions))
3 print('MSE:', metrics.mean_squared_error(y_test, predictions))
4 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

MAE: 25.590560026817133

MSE: 3777.721275095668

RMSE: 61.46317007034105


Exploratory Data Analysis

In [106]:

```
1 Raindf=df[['YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',
2           'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL']]
```


```
In [107]: 1 sns.heatmap(Raindf.corr())
```

```
Out[107]: <Axes: >
```




```
In [108]: 1 predictions=lm.predict(x_test)
          2 plt.scatter(y_test,predictions)
          3
```

Out[108]: <matplotlib.collections.PathCollection at 0x1f094e8e740>




```
In [119]: 1 sns.displot(x='YEAR',y='ANNUAL',data=df,color='aquamarine')  
          2 plt.show()
```




```
In [120]: 1 from sklearn.linear_model import Ridge,Lasso,RidgeCV
          2 ridgeReg = Ridge(alpha=10)
          3 ridgeReg.fit(x_train,y_train)
          4 train_score_ridge = ridgeReg.score(x_train,y_train)
          5 test_score_ridge = ridgeReg.score(x_test,y_test)
          6 print('\nRidge Model\n')
          7 print('Train score for ridge model is {}'.format(train_score_ridge))
          8 print('Test score for ridge model is {}'.format(test_score_ridge))
          9
```

Ridge Model

Train score for ridge model is 0.9916135364755825

Test score for ridge model is 0.9954207572738636

```
In [121]: 1 plt.figure(figsize = (10,4))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Ridge;$\alpha=10$')
3 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```



```
In [122]: 1 lassoReg = Lasso(alpha=10)
          2 lassoReg.fit(x_train,y_train)
          3 train_score_lasso = lassoReg.score(x_train,y_train)
          4 test_score_lasso = lassoReg.score(x_test,y_test)
          5
          6 print('\nRidge Model\n')
          7 print('Train score for lasso model is {}'.format(train_score_lasso))
          8 print('Test score for lasso model is{}'.format(test_score_lasso))
```


Ridge Model


Train score for lasso model is 0.9916132450486558

Test score for lasso model is0.9954172471874398

In [123]:

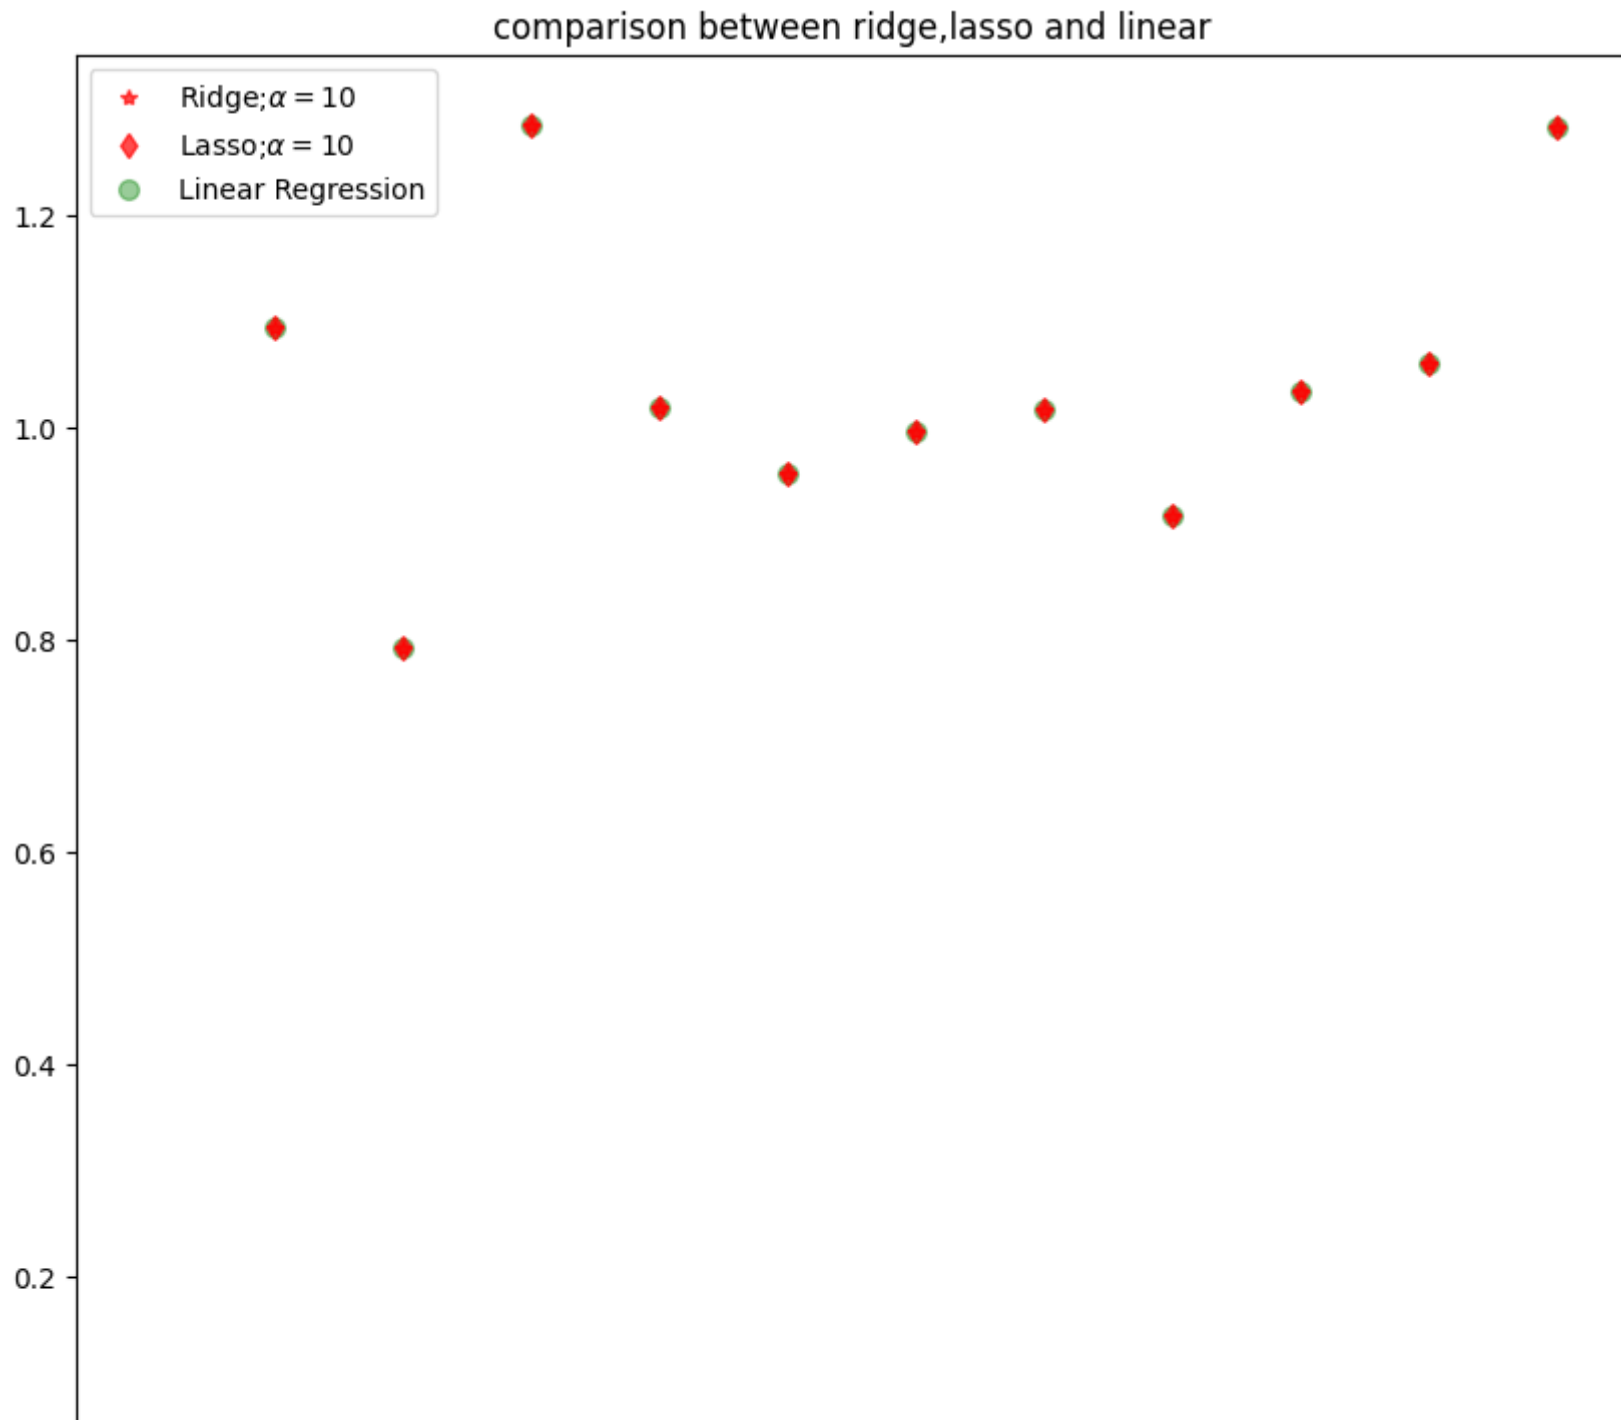
```
1 plt.figure(figsize = (10,10))
2 plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Lasso;$\lambda$')
3 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regres')
4 plt.xticks(rotation=90)
5 plt.legend()
6 plt.show()
```







In [124]:

```
1 plt.figure(figsize=(10,10))
2
3 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red',label=r'Ridge;$\
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='d',markersize=6,color='red',label=r'Lasso;$\
5 plt.plot(features,lm.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regres
6 plt.xticks(rotation=90)
7 plt.title('comparison between ridge,lasso and linear')
8 plt.legend()
9 plt.show()
```



```
In [125]: 1 # Linear CV model using Ridge
2 from sklearn.linear_model import RidgeCV
3 ridge_CV=RidgeCV(alphas=[0.0001,0.01,0.001,0.1]).fit(x_train,y_train)
4 print("The train score for ridge model is {}".format(ridge_CV.score(x_train,y_train)))
5 print("The test score for ridge model is {}".format(ridge_CV.score(x_test,y_test)))
```

The train score for ridge model is 0.9916135364727774

The test score for ridge model is 0.9954207508330051

```
In [126]: 1 # Linear CV model using Lasso
2 from sklearn.linear_model import LassoCV
3 lasso_CV=LassoCV(alphas=[0.0001,0.01,0.001,0.1]).fit(x_train,y_train)
4 print("The train score for lasso model is {}".format(lasso_CV.score(x_train,y_train)))
5 print("The test score for lasso model is {}".format(lasso_CV.score(x_test,y_test)))
```

The train score for lasso model is 0.9916135364460524

The test score for lasso model is 0.9954207194518694

```
In [127]: 1 from sklearn.linear_model import ElasticNet
          2 regr = ElasticNet()
          3 regr.fit(x,y)
          4 print(regr.coef_)
          5 print(regr.intercept_)
          6 y_pred_elastic = regr.predict(x_train)
          7 mean_squared_error = np.mean((y_pred_elastic-y_train)**2)
          8 print('Mean squared error on test set',mean_squared_error)
```

```
[0.00452878 1.21391301 1.00790437 1.04920427 0.96127995 1.00091858
 0.99101193 0.99152702 0.98668101 1.02070562 1.04790013 1.26485708]
-4.503727353915792
Mean squared error on test set 7054.045634195742
```

```
In [128]: 1 regr.score(x_train,y_train)
```

```
Out[128]: 0.9913834208596437
```

Conclusion

The accuracy is same for all Linear,Ridge,Lasso Regressions where Elastic Net has less accuracy compared to them.